

Homework 1 - Data Preparation Phase

Name/Znumber: Jordan Small / Z23465928 Professor : Juan Yepes Date : 31 JAN, 2024

Dataset: STUDENT PERFORMANCE DATASET

Description: This data is related to student achievement of a Portuguese school, obtained from using school reports and questionnaires. The classification goal is to predict student performance, particularly in a post-secondary math course.

<https://archive.ics.uci.edu/dataset/320/student+performance>

```
import pandas as pd #Import the pandas module
```

1. Load/read the csv file and display its shape

```
math = pd.read_csv('student-mat.csv', sep=';')  
#Check the shape of the object/dataframe  
math.shape  
(395, 33)
```

2. Print the header and the last few rows

```
#Print the header
```

```
math.head(n=5)
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
0	GP	F	18	U	GT3	A	4	4	at_home
1	GP	F	17	U	GT3	T	1	1	at_home
2	GP	F	15	U	LE3	T	1	1	at_home
3	GP	F	15	U	GT3	T	4	2	health
4	GP	F	16	U	GT3	T	3	3	other

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	6	5	6	6
1	5	3	3	1	1	3	4	5	5	6
2	4	3	2	2	3	3	10	7	8	10
3	3	2	2	1	1	5	2	15	14	15
4	4	3	2	1	2	5	4	6	10	10

```
[5 rows x 33 columns]
```

```
#Print the last few rows
```

```
math.tail(n=5)
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
Fjob \									
390	MS	M	20	U	LE3	A	2	2	services
services									
391	MS	M	17	U	LE3	T	3	1	services
services									
392	MS	M	21	R	GT3	T	1	1	other
other									
393	MS	M	18	R	LE3	T	3	2	services
other									
394	MS	M	19	U	LE3	T	1	1	other
at_home									

	...	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2
G3										
390	...	5	5	4	4	5	4	11	9	9
9										
391	...	2	4	5	3	4	2	3	14	16
16										
392	...	5	5	3	3	3	3	3	10	8
7										
393	...	4	4	1	3	4	5	0	11	12
10										
394	...	3	2	3	3	3	5	5	8	9
9										

```
[5 rows x 33 columns]
```

3. Print a summary of the dataset's statistical details

```
math.describe()
```

	age	Medu	Fedu	traveltime	studytime
failures \					
count	395.000000	395.000000	395.000000	395.000000	395.000000
395.000000					
mean	16.696203	2.749367	2.521519	1.448101	2.035443
0.334177					
std	1.276043	1.094735	1.088201	0.697505	0.839240
0.743651					
min	15.000000	0.000000	0.000000	1.000000	1.000000
0.000000					
25%	16.000000	2.000000	2.000000	1.000000	1.000000

```

0.000000
50%      17.000000      3.000000      2.000000      1.000000      2.000000
0.000000
75%      18.000000      4.000000      3.000000      2.000000      2.000000
0.000000
max      22.000000      4.000000      4.000000      4.000000      4.000000
3.000000

      famrel      freetime      goout      Dalc      Walc
health \
count 395.000000 395.000000 395.000000 395.000000 395.000000
395.000000
mean   3.944304   3.235443   3.108861   1.481013   2.291139
3.554430
std    0.896659   0.998862   1.113278   0.890741   1.287897
1.390303
min    1.000000   1.000000   1.000000   1.000000   1.000000
1.000000
25%    4.000000   3.000000   2.000000   1.000000   1.000000
3.000000
50%    4.000000   3.000000   3.000000   1.000000   2.000000
4.000000
75%    5.000000   4.000000   4.000000   2.000000   3.000000
5.000000
max    5.000000   5.000000   5.000000   5.000000   5.000000
5.000000

      absences      G1      G2      G3
count 395.000000 395.000000 395.000000 395.000000
mean   5.708861  10.908861  10.713924  10.415190
std    8.003096   3.319195   3.761505   4.581443
min    0.000000   3.000000   0.000000   0.000000
25%    0.000000   8.000000   9.000000   8.000000
50%    4.000000  11.000000  11.000000  11.000000
75%    8.000000  13.000000  13.000000  14.000000
max   75.000000  19.000000  19.000000  20.000000

```

4. Display a concise summary of the dataframe

```

math.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
#   Column      Non-Null Count  Dtype
---  -
0   school      395 non-null    object
1   sex          395 non-null    object
2   age          395 non-null    int64
3   address      395 non-null    object

```

4	famsize	395	non-null	object
5	Pstatus	395	non-null	object
6	Medu	395	non-null	int64
7	Fedu	395	non-null	int64
8	Mjob	395	non-null	object
9	Fjob	395	non-null	object
10	reason	395	non-null	object
11	guardian	395	non-null	object
12	traveltime	395	non-null	int64
13	studytime	395	non-null	int64
14	failures	395	non-null	int64
15	schoolsup	395	non-null	object
16	famsup	395	non-null	object
17	paid	395	non-null	object
18	activities	395	non-null	object
19	nursery	395	non-null	object
20	higher	395	non-null	object
21	internet	395	non-null	object
22	romantic	395	non-null	object
23	famrel	395	non-null	int64
24	freetime	395	non-null	int64
25	goout	395	non-null	int64
26	Dalc	395	non-null	int64
27	Walc	395	non-null	int64
28	health	395	non-null	int64
29	absences	395	non-null	int64
30	G1	395	non-null	int64
31	G2	395	non-null	int64
32	G3	395	non-null	int64

dtypes: int64(16), object(17)

memory usage: 102.0+ KB

5. Add an index column and display its new shape

```
math['index'] = pd.Series(range(0,395))
```

```
math.shape
```

```
(395, 34)
```

6. Choose a numerical field and display its unique values

```
print(math['absences'].unique())
```


```
[ 6  4 10  2  0 16 14  7  8 25 12 54 18 26 20 56 24 28  5 13 15 22  3
21
 1 75 30 19  9 11 38 40 23 17]
```

7. Replace an extreme set of values in the dataset with zero, -1, or np.NaN

```
#After searching through all numerical data columns, no extreme outliers  
#were found; instead, I replaced the absence value of 0, as seen below.  
  
import numpy as np  
np.NaN #Not a number  
  
nan  
  
#Replace all absences = 0 for np.NaN  
  
math['absences'] = math['absences'].replace({0: np.NaN})  
  
print(math['absences'].unique())  
  
[ 6.  4. 10.  2. nan 16. 14.  7.  8. 25. 12. 54. 18. 26. 20. 56. 24.  
28.  
 5. 13. 15. 22.  3. 21.  1. 75. 30. 19.  9. 11. 38. 40. 23. 17.]
```

8. Plot a histogram of a numerical value

```
math['age'].plot(kind = 'hist', title = 'Histogram of Student Ages')  
  
<Axes: title={'center': 'Histogram of Student Ages'},  
ylabel='Frequency'>
```



9. Choose a categorical field and represent it as a numerical field (create a new column)

```
math["Pstatus_numeric"] = math['Pstatus'] #Replicate the Pstatus variable
```

```
print(math['Pstatus'].unique()) #Print/inspect unique values for Pstatus
```

```
['A' 'T']
```

```
#Create a dictionary with the numeric value equivalencies
```

```
dict_Pstatus = {"Pstatus_numeric": {'A':0,'T':1}}
dict_Pstatus
```

```
{'Pstatus_numeric': {'A': 0, 'T': 1}}
```

```
#Make the replacement and display the last few rows
```

```
math.replace(dict_Pstatus, inplace=True)
```

```
math.tail(n=5)
```

```

  school sex  age address famsize Pstatus  Medu  Fedu    Mjob
Fjob  \

```

```

390      MS  M  20      U  LE3      A  2  2  services
services
391      MS  M  17      U  LE3      T  3  1  services
services
392      MS  M  21      R  GT3      T  1  1  other
other
393      MS  M  18      R  LE3      T  3  2  services
other
394      MS  M  19      U  LE3      T  1  1  other
at_home

      ...  Dalc  Walc  health  absences  G1  G2  G3  index  Pstatus_numeric
\
390  ...    4    5      4      11    9  9  9  390      0
391  ...    3    4      2      3   14 16 16  391      1
392  ...    3    3      3      3   10  8  7  392      1
393  ...    3    4      5      0   11 12 10  393      1
394  ...    3    3      5      5    8  9  9  394      1

      age_z
390  2.592380
391  0.238380
392  3.377047
393  1.023046
394  1.807713

[5 rows x 36 columns]

```

10. Use the zscore function to standardize a numerical field (create a new column)

$$z = \frac{\text{Standardized Value}}{\text{Standard deviation}} = \frac{x - \bar{x}}{s} = \frac{\text{Data value} - \text{mean}}{\text{Standard deviation}}$$

```

#Import the scipy module

from scipy import stats

#Number of standard deviations above the mean

math["age_z"] = stats.zscore(math['age'])
print(math['age_z'].head(n=20))

0      1.023046
1      0.238380


```

```
2    -1.330954
3    -1.330954
4    -0.546287
5    -0.546287
6    -0.546287
7     0.238380
8    -1.330954
9    -1.330954
10   -1.330954
11   -1.330954
12   -1.330954
13   -1.330954
14   -1.330954
15   -0.546287
16   -0.546287
17   -0.546287
18    0.238380
19   -0.546287
Name: age_z, dtype: float64
```

```
#Age Histogram (not normalized)
```

```
math['age'].plot(kind = 'hist', title = 'Histogram of Age')
```

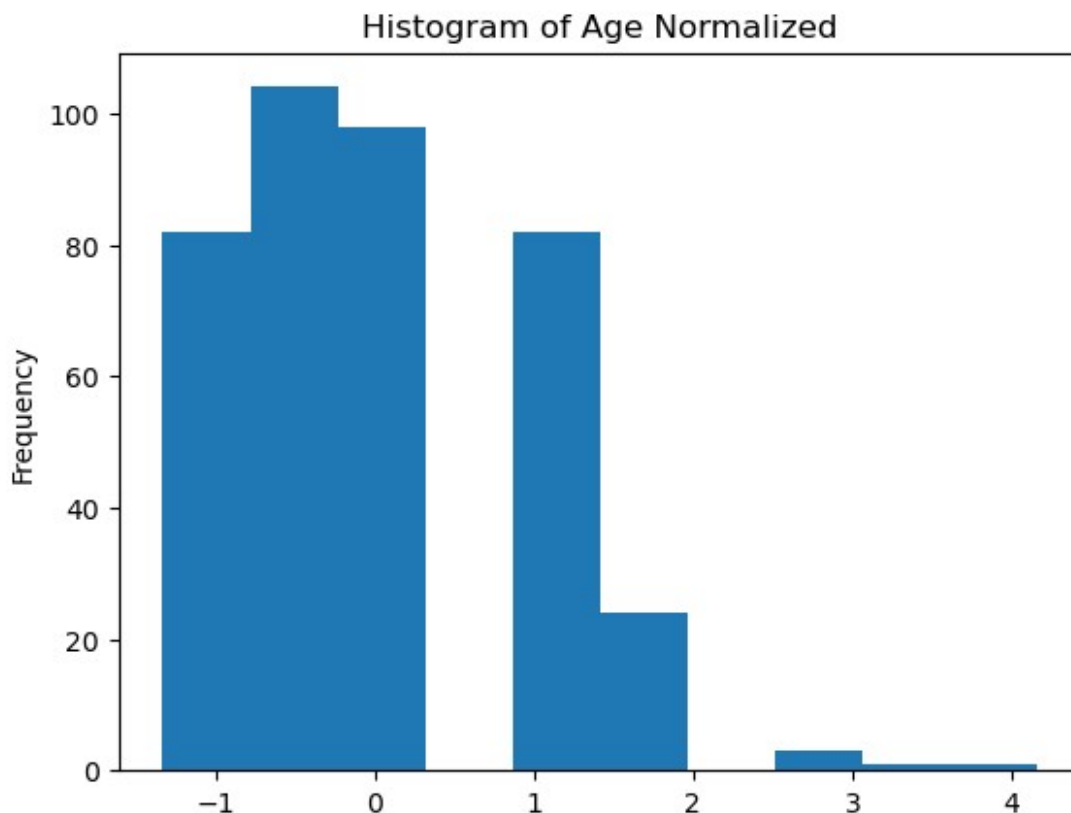
```
<Axes: title={'center': 'Histogram of Age'}, ylabel='Frequency'>
```

```
#Age Histogram (Normalized)
```

```
math['age_z'].plot(kind = 'hist', title = 'Histogram of Age  
Normalized')
```

```
<Axes: title={'center': 'Histogram of Age Normalized'},  
ylabel='Frequency'>
```



11. Identify a field and use a criterion of your choosing to filter for outliers. Create a new dataset with the outliers

```
math.query('age_z > 3 | age_z < -3') #List the outliers
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
Fjob \									
247	GP	M	22	U	GT3	T	3	1	services
392	MS	M	21	R	GT3	T	1	1	other

	...	Dalc	Walc	health	absences	G1	G2	G3	index	Pstatus_numeric
age_z										
247	...	5	5	1	16	6	8	8	247	1
392	...	3	3	3	3	10	8	7	392	1

[2 rows x 36 columns]

```
#Create a new dataset with the outlier values
```

```
math_outliers = math.query('age_z > 3 | age_z < -3')
```

#Display the first few rows of outliers

```
math_outliers.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob
Fjob \									
247	GP	M	22	U	GT3	T	3	1	services
392	MS	M	21	R	GT3	T	1	1	other

	...	Dalc	Walc	health	absences	G1	G2	G3	index	Pstatus_numeric
age_z										
247	...	5	5	1	16	6	8	8	247	1
392	...	3	3	3	3	10	8	7	392	1

[2 rows x 36 columns]

12. Sort the dataset and display 15 interesting fields

#Sort the dataset

```
math_sort = math.sort_values(['age_z'], ascending = False)
```

#Display 15 interesting fields

```
print(math_sort[['age', 'failures']].head(n=15))
```

	age	failures
247	22	3
392	21	3
390	20	2
306	20	0
376	20	2
394	19	0
310	19	1
340	19	1
336	19	1
315	19	1
314	19	2
313	19	1
127	19	3
311	19	0
307	19	1